

Canvas 1.4 Python API

Introduction

This document describes the Python application programming interface that accompanies Canvas version 1.4, which is part of the Schrödinger Suite 2011 release.

The Canvas Python API consists of a rich library of object-oriented tools for developing custom cheminformatics applications with a range of functionality, including:

- Chemical structure storage and I/O
- Substructure matching
- Fingerprinting creation and I/O
- Fingerprint bit mapping
- Similarity/distance calculations
- Hierarchical clustering

Each Python application you create must import the `canvas` module:

```
import schrodinger.application.canvas.base as canvas
```

If you wish to use `mmct` objects and convert back and forth between Canvas `ChmMol` objects, you will need additional imports:

```
from schrodinger import structure
from schrodinger.infra.canvas import ChmMmctAdaptor
```

Furthermore, you must have a valid `CANVAS_FULL` license, and the application must successfully checkout that license before the Canvas APIs can be used. The following checkout procedure is recommended:

```
try:
    lic = canvas.ChmLicenseFull()
except:
    sys.stderr.write("Unable to checkout CANVAS_FULL license\n")
    sys.exit(1)

if not lic.isValid():
    sys.stderr.write("CANVAS_FULL license is not valid\n")
    sys.exit(1)
```

To ensure that all environment variables and paths are set correctly, a Canvas Python application should be invoked as follows:

```
$SCHRODINGER/run <app_name> [options]
```

Where `$SCHRODINGER` points to the directory in which the 2011 Schrodinger software is installed.

Constructors and Static Methods

Most Canvas objects are created automatically as a returned value from a member function invoked on another object. For example, if `mol` is a `ChmMol` object, `mol.getAtom(i)` returns a `ChmAtom` object. However, there are certain objects that you will create explicitly through a constructor, for example,

```
query = canvas.ChmQuery("c1nc[c,n]cc1")

bitset = canvas.ChmBitset(1024)
```

There are also cases where you will create an object through a static method of that class, such as,

```
mol = canvas.ChmMol.fromSMILES("Nc1cccccl")

query = canvas.ChmQuery.fromMDLFile("query.mol")
```

In the material that follows, any applicable constructors are documented before other methods of the class, and all static methods are noted as such.

Iterable Lists

Lists returned by Canvas APIs behave much like ordinary Python lists, but they are not entirely equivalent, nor can you supply an ordinary Python list as an argument in an API that calls for a Canvas-style list. Another important difference is that after iterating through a Canvas list, you must rewind it before you can iterate through it again. For example:

```
atoms = mol.getAtoms()

print "First pass... "
for atom in atoms:
    print "Label =",atom.getLabel()

print "Second pass... "
atoms.rewind()
for atom in atoms:
    print "Label =",atom.getLabel()
```

Note that the “for” loops above are equivalent to the following:

```
while atoms.hasNext():
    atom = atoms.next()
    print "Label =",atom.getLabel()
```

Methods

[*i*]

The i^{th} item in the list, where i runs from 0 to `size()-1`. The legality of i is not checked, so if the supplied value is out of range, a segmentation fault can occur.

`at(i)`

The i^{th} item in the list, where i runs from 0 to `size()-1`. If an illegal value is supplied, an exception is thrown.

`hasNext()`

True if more items remain in the list.

`next()`

The next item in the list

`rewind()`

Rewinds to the beginning of the list.

`size()`

The total number of items in the list.

Input/Output Streams

Certain Canvas APIs operate on an input or output stream that's attached to file. If the input and output file names are *inFile* and *outFile*, respectively, the input/output stream objects may be created as follows:

```
inStream = canvas.istream(inFile)
outStream = canvas.ofstream(outFile)
```

Exceptions

When an error occurs within a Canvas constructor or method, a `RuntimeError` exception with an informative error message is normally thrown. To handle the exception and display the error message, the following approach may be used:

```
try:
    <Call some Canvas function>
except Exception as canvasErr:
    sys.stderr.write(str(canvasErr) + "\n")
```

Chemical Structure APIs

The classes and methods in this section support chemical structure storage and I/O.

ChmAtom – Describes a single atom in a ChmMol.

Methods

findNeighbor(*otherAtom*)

Loops over the bonds to the current ChmAtom, and returns the position (0, 1,...) at which the bond to the supplied ChmAtom is found.

firstNeighbor()

The first ChmAtom bonded to the current ChmAtom.

getAtomicNumber()

Atomic number.

getBond(*otherAtom*)

The ChmBond between the current ChmAtom and the supplied ChmAtom.

getBond(*i*)

The ChmBond between the current ChmAtom and the i^{th} ChmAtom to which it is bonded. All atoms (explicit + implicit) are considered, so i ranges from 0 to getBondCount(True)-1.

getBondCount([*wantImplicitHydrogens*])

Number of atoms bonded to the current ChmAtom. The Boolean *wantImplicitHydrogens* is False by default and determines whether bonds to implicit hydrogens should be counted.

getBonds([*wantImplicitHydrogens*])

Iterable list of bonds (as ChmBond objects) made by the current ChmAtom. The Boolean *wantImplicitHydrogens* is False by default and determines whether bonds to implicit hydrogens will not be included.

getFormalCharge()

Integer formal charge.

getHeavyBondCount()

Number of heavy atoms bonded to the current ChmAtom.

getHeavyBonds()

Iterable list of bonds (as ChmBond objects) made to heavy atoms.

getHeavyNeighbors()

Iterable list of heavy atoms (as ChmAtom objects) bonded to the current ChmAtom.

getHybridization()

Hybridization as an integer: 0→unknown, 1→sp, 2→sp², 3→sp³, 4→sp³d, 5→sp³d².

getHydrogenCount()

Number of hydrogens (explicit + implicit) bonded to the current ChmAtom.

getImplicitHydrogenCount()

Number of implicit hydrogens bonded to the current ChmAtom.

getImplicitNeighbors()

Iterable list of bonds (as ChmBond objects) to implicit hydrogens.

getLabel()

A string that contains a label of the form <symbol><number>, where <symbol> is the elemental symbol, and <number> is the zero-based atom number (0, 1,...).

getMolIndex()

Index (0, 1,...) of the current ChmAtom in its parent ChmMol.

getNeighbors([*wantImplicitHydrogens*[, *wantExplicitHydrogens*]])

Iterable list of atoms (as ChmAtom objects) bonded to the current ChmAtom. The Boolean *wantImplicitHydrogens* is False by default and determines whether implicit hydrogens should be included. The Boolean *wantExplicitHydrogens* is True by default and determines whether explicit hydrogens should be included.

getRingSize()

Size of the smallest ring that contains the current ChmAtom.

getSymbol()

Elemental symbol.

getTotalValence()

Total valence, which is the sum of the bond orders over all bonds made by the current ChmAtom. For this computation, single→1, double→2, triple→3, aromatic→1.5.

getUserData()

An integer-valued piece of data associated with the atom.

getX()

X coordinate.

getY()

Y coordinate.

getZ()

Z coordinate.

hasCoords()
True if the ChmAtom has coordinates.

isAromatic()
True if in aromatic ring.

isBondedTo(*otherAtom*)
True if bonded to the supplied ChmAtom.

isCarbon()
True if carbon.

isHalogen()
True if F, Cl, Br or I.

isHeavy()
True if not hydrogen.

isHydrogen()
True if hydrogen.

isInRing()
True if in ring.

isNitrogen()
True if nitrogen.

isOxygen()
True if oxygen.

isSulfur()
True if sulfur.

isTerminal()
True if terminal.

setPartialCharge(*q*)
Sets the partial charge.

setUserData(*value*)
Sets an integer-valued piece of data to associate with the atom.

setX(*x*)
Sets the X coordinate.

setY(y)

Sets the Y coordinate.

setZ(z)

Sets the Z coordinate.

ChmBond – Describes a bond between two ChmAtom objects.

Methods

atom1()

First ChmAtom in the ChmBond.

atom2()

Second ChmAtom in the ChmBond.

getBondLength()

Bond length in angstroms.

getMolIndex()

Index (0, 1,...) of the current ChmBond in its parent ChmMol.

getOrder()

Bond order. 0→unknown, 1→single, 2→double, 3→triple, 4→aromatic, 5→Kekulized single, 6→Kekulized double.

getRingSize()

Size of the smallest ring that contains the current ChmBond.

getUserCode()

A non-negative integer-valued code associated with the bond.

isAromatic()

True if in aromatic ring.

isInRing()

True if in ring.

isTerminal()

True if terminal.

mate(*atom*)

Given one ChmAtom in the current ChmBond, this function returns the other ChmAtom.

setUserCode(*code*)

Sets a non-negative integer-valued code to associate with the bond.

ChmMmctAdaptor – Converts between ChmMol and mmct objects.

Stereochemistry Constants

ChmMmctAdaptor.NoStereo – Ignore stereochemistry (0).

ChmMmctAdaptor.StereoFromGeometry – Use mmstereo library to assign stereochemistry from 3D structure (1).

ChmMmctAdaptor.StereoFromGeometry_Safe – As above, but ignore mmstereo assignments that Canvas doesn't agree with (2).

ChmMmctAdaptor.StereoFromAnnotation – Use existing mmstero annotations (3).

ChmMmctAdaptor.StereoFromAnnotation_Safe – As above, but ignore annotations that Canvas doesn't agree with (4).

ChmMmctAdaptor.StereoFromAnnotationAndGeometry – Use existing mmstero annotations, with assignment from 3D structure as a backup for unspecified stereochemistry (5).

ChmMmctAdaptor.StereoFromAnnotationAndGeometry_Safe – As above, but ignore annotations/assignments that Canvas doesn't agree with (6).

ChmMmctAdaptor.StereoFromCurrentSettings – Use mmstereo library to assign stereochemistry from 3D structure. Existing annotations are used or ignored, depending on the current mmstereo settings (7).

ChmMmctAdaptor.StereoFromCurrentSettings_Safe – As above, but ignore assignments/annotations that Canvas doesn't agree with (8).

Methods

`create(ctHandle, [stereoTreatment[, allHydrogensAsImplicit[, wantProperties]])`

Does mmct→ChmMol conversion and returns a ChmMol object. The integer *stereoTreatment* is ChmMmct.NoStereo by default and controls how stereochemistry is transferred. See “Stereochemistry Constants” for other legal values. The Boolean *allHydrogensAsImplicit* is False by default and determines whether all hydrogens should be stored as implicit hydrogens. The Boolean *wantProperties* is True by default and determines whether properties should be copied.

`create(mol)`

Does ChmMol→mmct conversion and returns an mmct handle.

ChmMol – Describes a chemical structure.

Methods

`fromMDL(molString[, autoName])`

Static method that creates a ChmMol from a MOL file string or a single-structure SD file string. The Boolean *autoName* is False by default and determines whether a name is generated automatically with each call to this function (“Molecule1”, “Molecule2”, etc.).

`fromSMILES(smiles[, autoName])`

Static method that creates a ChmMol from a SMILES string. The Boolean *autoName* is False by default and determines whether a name is generated automatically with each call to this function (“Molecule1”, “Molecule2”, etc.).

`getAtom(i)`

The i^{th} ChmAtom in the ChmMol. All atoms (explicit and implicit) are considered, so *i* ranges from 0 to `getAtomCount(True)-1`.

`getAtomCount([wantImplicitHydrogens])`

Number of atoms in the ChmMol. The Boolean *wantImplicitHydrogens* is False by default and determines whether implicit hydrogens should be counted.

`getAtoms([wantImplicitHydrogens])`

Iterable list of atoms (as ChmAtom objects) in the ChmMol. The Boolean *wantImplicitHydrogens* is False by default and determines whether implicit hydrogens should be included.

`getAtoms(mask)`

Iterable list of a subset of atoms in the ChmMol. The argument *mask* is a ChmBitset object (see Bitset APIs), with “on” bits for the subset of atoms in question. The length of *mask* should be `getAtomCount(False)`.

`getBond(i)`

The i^{th} ChmBond in the ChmMol. All atoms (explicit and implicit) are considered, so *i* ranges from 0 to `getBondCount(True)-1`.

`getBondCount([wantImplicitHydrogens])`

Number of bonds in the ChmMol. The Boolean *wantImplicitHydrogens* is False by default and determines whether bonds to implicit hydrogens should be counted.

`getBonds([wantImplicitHydrogens])`

Iterable list of bonds (as ChmBond objects) in the ChmMol. The Boolean *wantImplicitHydrogens* is False by default and determines whether bonds to implicit hydrogens should be included.

`getBonds(bondMask)`

Iterable list of a subset of bonds in the ChmMol. The argument *bondMask* is a ChmBitset object (see Bitset APIs), with “on” bits for the subset of bonds in question.

`getBondsForAtoms(atomMask)`

Iterable list of bonds that connect a subset of atoms. The argument *atomMask* is a ChmBitset with “on” bits for a subset of atoms in the structure. The atoms need not form a contiguous substructure.

`getConnectionTable([wantImplicitHydrogens])`

Connection table: 1→bonded, 0→not bonded. The Boolean *wantImplicitHydrogens* is False by default and determines whether implicit hydrogens should be included. Use `<object>.getItem(i, j)` on the returned object to get the connectivity between atoms *i* and *j*. Atom numbers *i* and *j* run from 0 to `getAtomCount(wantImplicitHydrogens)-1`.

`getConnectionTableWithBondOrders([wantImplicitHydrogens])`

Analogous to `getConnectionTable`, except that `<object>.getItem(i, j)` returns the bond order if atoms *i* and *j* are bonded.

`getDistanceMatrix([wantImplicitHydrogens])`

Shortest path distance matrix. The Boolean *wantImplicitHydrogens* is False by default and determines whether implicit hydrogens should be included. Use `<object>.getItem(i, j)` on the returned object to get the distance between atoms *i* and *j*. Atom numbers *i* and *j* run from 0 to `getAtomCount(wantImplicitHydrogens)-1`.

`getDistanceMatrix3D([wantImplicitHydrogens], excludeZ])`

Three-dimensional distance matrix. The Boolean *wantImplicitHydrogens* is False by default and determines whether implicit hydrogens should be included. The Boolean *excludeZ* is False by default and determines whether Z coordinates should be ignored when computing distances. Use `<object>.getItem(i, j)` on the returned object to get the distance between atoms *i* and *j*. Atom numbers *i* and *j* run from 0 to `getAtomCount(wantImplicitHydrogens)-1`.

`getHeavyAtomCount()`

Number of heavy atoms.

`getMW()`

Molecular weight.

`getName()`

Structure name.

`getNetCharge()`

Net charge on structure.

`getProperties()`

Name→value mappings of all stored properties. May be coerced into an ordinary Python dictionary as follows:

```
props = dict(mol.getProperties())
```

`getProperty(propName)`

The value of a stored property, cast as a string.

`getPropertyCount()`

The number of stored properties.

`has2DCoords()`

True if structure has X and Y coordinates only.

`has3DCoords()`

True if structure has X, Y and Z coordinates.

`hasCoords()`

True if structure has 2D or 3D coordinates.

`hasProperty(propertyName)`

True if *propertyName* is one of the stored properties.

`setName(name)`

Sets structure name.

`setProperty(propertyName, propertyValue)`

Sets/adds a property, where *propertyValue* is a string.

`toCanonicalSMILES()`

Creates a canonical SMILES string from the structure.

`toMDL([withHydrogens[], withProps[], withCoords[], name])`

Creates an SD file string from the structure. The Boolean *withHydrogens* is False by default and determines whether hydrogens should be included. The Boolean *withProps* is True by default and determines whether stored properties should be output. The Boolean *withCoords* is True by default and determines whether the existing coordinates should be retained (zeros are written if False). The string *name* is used to override automatic name assignment (“Molecule1”, “Molecule2”, etc.).

`toSMILES([wantKekulized])`

Creates an ordinary SMILES string from the structure. The Boolean *wantKekulized* is False by default and determines whether aromatic rings should be represented using alternating single/double bonds.

ChmSDReader – Iterable SD file reader. Note that the following two methods for reading a file are equivalent:

```
1. sdr = canvas.ChmSDReader("file.sdf")
   for mol in sdr:
       print "Structure name =",mol.getName()

2. sdr = canvas.ChmSDReader("file.sdf")
   while sdr.hasNext():
       mol = sdr._next()
       print "Structure name =",mol.getName()
```

Constructors

ChmSDReader(*sdFile*)

Creates an SD file reader object and opens it on the specified file.

Methods

_next()

The next ChmMol in the file.

getError()

Description of most recent error.

hasError()

True if the most recent attempt to read produced an error.

hasNext()

True if more structures remain in the file.

ChmSmilesReader – Reads a file containing SMILES strings, one per line. If a SMILES string is followed by a space and a second field, that field will be stored as the structure name. A typical workflow to read a file is as follows:

```
smi = canvas.ChmSmilesFileReader("file.smi")
while smi.hasNextMol():
    mol = smi.nextMol()
    print "Structure name =",mol.getName()
```

Constructors

ChmSmilesFileReader(*smilesFile*)

Creates a SMILES file reader object and opens it on the specified file.

Methods

getError()

Description of most recent error.

getFailureCount()

Total number of errors.

getLastName()

The most recent structure name. If no name was found in the file, it's generated automatically as Molecule N , where N is the structure count (1, 2,...).

getLastSMILES()

The most recent SMILES string read.

getSuccessCount()

Total number of SMILES successfully read.

hasError()

True if the most recent attempt to read produced an error.

hasNextMol()

True if more SMILES strings remain in the file.

nextMol()

ChmMol generated from the next SMILES string in the file.

skip()

Skip over the next SMILES string.

Substructure Matching APIs

The classes and methods in this section support the creation and execution of 2D chemical queries.

ChmMatch – Describes a 2D match.

Methods

getMatchedAtoms()

Iterable list of the atoms (as ChmAtom objects) matched by the query.

ChmQuery – A 2D chemical query. Used in conjunction with ChmQueryMatcher.

Constructors

ChmQuery(*smarts*)

Creates a ChmQuery from a SMARTS string.

ChmQuery(*mol*)

Creates a ChmQuery from a ChmMol.

Methods

getAtomCount()

The number of atoms in the ChmQuery.

getSource()

The SMARTS representation of the query. This is a valid method even if the ChmQuery was created from a ChmMol.

fromMDL(*molString*)

Static method that creates a ChmQuery from a MOL file string or a single-structure SD file string.

fromMDLFile(*fileName*)

Static method that creates a ChmQuery from a MOL file or a single-structure SD file.

ChmQueryMatcher. Matches a ChmQuery to a ChmMol.

Constructors

ChmQueryMatcher([*uniqueFilter*])

The Boolean *uniqueFilter* is True by default and determines whether only a single mapping should be returned when the query can be mapped to a set of atoms in more than one way (e.g., the query c1ccccc1 can be mapped to a phenyl ring in 12 ways).

Methods

getUniqueFilter()

True if only one mapping per match should be returned.

hasExactMatch(*query*, *mol*)

True if the supplied ChmQuery produces an exact match to the supplied ChmMol.

hasMatch(*query*, *mol*)

True if the supplied ChmQuery produces at least one match to the supplied ChmMol.

match(*query*, *mol*)

Matches a ChmQuery to a ChmMol and returns an iterable list of ChmMatch objects.

`matchCount(query, mol)`

The number of times a ChmQuery matches a ChmMol.

`matchMask(query, mol)`

Matches a ChmQuery to a ChmMol and returns an iterable list of matches as ChmBitset objects (see Bitset APIs).

`matchMaskSubset(query, mol, bitset)`

Matches a ChmQuery to a subset of atoms in a ChmMol. The argument *bitset* is a ChmBitset object (see Bitset APIs), with “on” bits for the subset of atoms in question. The length of *bitset* must be `mol.getAtomCount(False)` or `mol.getAtomCount(True)`.

`matchSubset(query, mol, atoms)`

Matches a ChmQuery to a subset of atoms in a ChmMol. The argument *atoms* is a list of ChmAtom objects returned by another Canvas API, such as `mol.getAtoms(bitset)`. It cannot be an ordinary Python list.

`setUniqueFilter(uniqueFilter)`

Turns unique filtering on/off according to the value of the Boolean *uniqueFilter*.

Bitset APIs

The classes and methods in this section support the storage, manipulation and comparison of fingerprint data. See “Fingerprint APIs” for the actual creation of fingerprints from chemical structures.

ChmBitset. A fixed length bit string with explicit on/off values (compare ChmSparseBitset). For the sake of clarity in describing certain APIs, the example bitset “1,0,0,1,1” will be used. This bitset has a length of 5 with “on” bits at positions 0, 3, and 4.

Similarity/Distance Measures

A large number of methods are available for computing the similarity or distance between two ChmBitset objects *bitset1* and *bitset2*. Some of these methods will not be familiar to most users, so precise definitions are supplied here.

Let:

$a \equiv$ Count of “on” bits in *bitset1*.

$b \equiv$ Count of “on” bits in *bitset2*.

$c \equiv$ Count of bits that are “on” in both *bitset1* and *bitset2*.

$d \equiv$ Count of bits that are “off” in both *bitset1* and *bitset2*.

$A \equiv a - c =$ count of bits that are exclusively “on” bits in *bitset1*.

$B \equiv b - c =$ count of bits that are exclusively “on” bits in *bitset2*.

$s \equiv$ length of bitset.

Buser similarity	$(\sqrt{c \cdot d} + c) / (\sqrt{c \cdot d} + a + b - c)$
Cosine similarity	$c / \sqrt{a \cdot b}$
Dice similarity	$c / (0.5 \cdot (a + b))$
Dixon distance	$(1.0 - \text{Tanimoto}) \cdot \text{Hamming}$
Euclidean distance	$\sqrt{A + B}$
Hamann similarity	$(c + d - A - B) / (A + B + c + d)$
Hamming distance	$A + B$
Kulczynski similarity	$0.5 \cdot (c/a + c/b)$
Matching similarity	$(c + d) / s$
McConnaughey similarity	$((c \cdot c) - ((a - c) \cdot (b - c))) / (a \cdot b)$
MinMax similarity	$\sum_i \min[\text{bitset1}(i), \text{bitset2}(i)] / \max[\text{bitset1}(i), \text{bitset2}(i)]$
Modified Tanimoto similarity	$\alpha \text{Tanimoto} + (1 - \alpha) \cdot T_0$ where, $\alpha \equiv 2/3 - (a + b) / [6 \cdot \min(d, 10000)]$ $T_0 \equiv d / (a + b - 2 \cdot c + d) = \text{Tanimoto of "off" bits}$
Pattern difference distance	$(A \cdot B) / ((A + B + c + d)^2)$
Pearson similarity	$((c \cdot d) - (A \cdot B)) / \sqrt{a \cdot b \cdot (A + d) \cdot (B + d)}$
Petke similarity	$c / \max(a, b)$
Rogers Tanimoto similarity	$(c + d) / (2 \cdot (a + b) - (3 \cdot c) + d)$
Shape distance	$(A - B) / (A + B + c + d) - ((A - B) / (A + B + c + d))^2$
Simpson similarity	$c / \min(a, b)$
Size distance	$(A - B)^2 / (A + B + c + d)^2$
Soergel distance	$(A + B) / (A + B + c)$
Tanimoto similarity	$c / (a + b - c)$
Tversky(α, β) similarity	$c / (\alpha(a - c) + \beta(b - c) + c)$
Variance distance	$(A + B) / (4 \cdot (A + B + c + d))$
Yule similarity	$(c \cdot d - A \cdot B) / (c \cdot d + A \cdot B)$

Constructors

ChmBitset(*n*)

Creates a bitset of length *n* with all bits off. The length automatically increases if bits higher than *n* - 1 are turned on.

ChmBitset(*onBits*, *n*)

Creates a bitset of length *n* with a specific collection of bits turned on. The argument *onBits* is an ordinary Python list or tuple, containing the positions of the on bits. The example bitset "1,0,0,1,1" could be created using either of the following statements:

```
bitset = canvas.ChmBitset([0, 3, 4], 5)
bitset = canvas.ChmBitset((0, 3, 4), 5)
```

The length automatically increases if bits higher than *n* - 1 are turned on, including bits in *onBits*.

Methods

`all()`

True if all bits are on.

`any()`

True if any bits are on.

`clear()`

Turns off all bits.

`clear(i)`

Turns off bit *i*. *i* runs from 0 to `size()-1`.

`compare(bitset)`

Compares the current ChmBitset to the supplied one in a bitwise numeric fashion. Returns 1, 0, or -1, which correspond, respectively, to the current bitset being greater than, equal to, or less than the supplied bitset.

`contains(bitset)`

True if all the on bits in the supplied ChmBitset are also on in the current ChmBitset. The two bitsets must have the same length.

`count()`

The number of on bits.

`countCommonOff(bitset)`

The number of off bits shared by the current ChmBitset and the supplied ChmBitset. The two bitsets must have the same length.

`countCommonOn(bitset)`

The number of on bits shared by the current ChmBitset and the supplied ChmBitset. The two bitsets must have the same length.

`countDifference(bitset)`

The number of bit differences between the current ChmBitset and the supplied ChmBitset. The two bitsets must have the same length.

`distDixon(bitset)`

Dixon distance between the current ChmBitset and the supplied one.

`distEuclidean(bitset)`

Euclidean distance.

`distHamming(bitset)`

Hamming distance.

`distPatternDifference(bitset)`
Pattern difference distance.

`distShape(bitset)`
Shape distance.

`distSize(bitset)`
Size distance.

`distSoergel(bitset)`
Soergel distance.

`distSquaredEuclidean(bitset)`
Squared Euclidean distance.

`distVariance(bitset)`
Variance distance.

`flip()`
Flips all bits.

`flip(i)`
Flips bit *i*. If $i \geq \text{size}()$, the length of the `bitset` is automatically increased.

`get(i)`
True if bit *i* is on. Returns False if *i* is larger than `size()-1`.

`hasDifference(bitset)`
True if any bit in the current `ChmBitset` differs from the corresponding bit in the supplied `ChmBitset`. The two bitsets must have the same length.

`isContainedBy(bitset)`
True if all the on bits in the current `ChmBitset` are also on in the supplied `ChmBitset`. The two bitsets must have the same length.

`intersects(bitset)`
True if the current `ChmBitset` and the supplied `ChmBitset` share at least one on bit. The two bitsets must have the same length.

`none()`
True if all bits are off.

`resize(n)`
Changes the length of the `bitset`. If shorter, bits are removed from the end; if longer, off bits are appended.

`set()`

Turns on all bits.

`set(i)`

Turns on bit *i*. If $i \geq \text{size}()$, the length of the bitset is automatically increased.

`simBuser(bitset)`

Buser similarity.

`simCosine(bitset)`

Cosine similarity.

`simDice(bitset)`

Dice similarity.

`simHamann(bitset)`

Hamann similarity.

`simKulczynski(bitset)`

Kulczynski similarity.

`simMatching(bitset)`

Matching similarity.

`simMcConnaughey(bitset)`

McConnaughey similarity.

`simMinMax(bitset)`

MinMax similarity.

`simModifiedTanimoto(bitset)`

Modified Tanimoto similarity.

`simPearson(bitset)`

Pearson similarity.

`simPetke(bitset)`

Petke similarity.

`simRogersTanimoto(bitset)`

Rogers Tanimoto.

`simSimpson(bitset)`

Simpson similarity.

`simTanimoto(bitset)`

Tanimoto similarity.

`simTversky(bitset, α , β)`

Tversky similarity.

`simYule(bitset)`

Yule similarity.

`size()`

The length of the bitset.

`to_bitstring()`

A comma-delimited string that represents all bit values. The example bitset would return “1,0,0,1,1”.

`toString()`

A comma-delimited string that contains the on bit positions. The example bitset would return “0,3,4”.

`toVector()`

An ordinary Python tuple of the on bit positions. The example bitset would return (0, 3, 4).

ChmSparseBitset. A bitset of theoretical length 2^{32} , which stores only positions of the “on” bits.

Similarity/Distance Measures

See ChmBitset for definitions.

Constructors

`ChmSparseBitset()`

Creates a sparse bitset with no on bits.

`ChmSparseBitset(onBits)`

Creates a sparse bitset with a specific collection of bits turned on. The argument *onBits* is an ordinary Python list or tuple, containing the positions of the on bits. The example bitset above could be created using either of the following statements:

```
bitset = canvas.ChmSparseBitset([0, 3, 4])
bitset = canvas.ChmSparseBitset((0, 3, 4))
```

Methods

`contains(bitset)`

True if all the on bits in the supplied ChmSparseBitset are also on in the current ChmSparseBitset.

count()

The number of on bits.

countCommonOff(*bitset*)

The number of off bits shared by the current ChmSparseBitset and the supplied ChmSparseBitset. This is normally a very large number since only a very small fraction of the 2^{32} bits in a ChmSparseBitset are typically on.

countCommonOn(*bitset*)

The number of on bits shared by the current ChmSparseBitset and the supplied ChmSparseBitset.

countDifference(*bitset*)

The number of bit differences between the current ChmSparseBitset and the supplied ChmSparseBitset.

density()

The fraction of bits that are on, i.e., count()/size().

distDixon(*bitset*)

Dixon distance between the current ChmSparseBitset and the supplied one.

distEuclidean(*bitset*)

Euclidean distance.

distHamming(*bitset*)

Hamming distance.

distPatternDifference(*bitset*)

Pattern difference distance.

distShape(*bitset*)

Shape distance.

distSize(*bitset*)

Size distance.

distSoergel(*bitset*)

Soergel distance.

distSquaredEuclidean(*bitset*)

Squared Euclidean distance.

empty()

True if no bits are on.

hasDifference(*bitset*)

True any bit in the current ChmSparseBitset differs from the corresponding bit in the supplied ChmSparseBitset.

isSubsetOf(*bitset*)

True if all the on bits in the current ChmSparseBitset are also on in the supplied ChmSparseBitset.

intersects(*bitset*)

True if the current ChmSparseBitset and the supplied ChmSparseBitset share at least one on bit.

positionOf(*i*)

If bit *i* is on, its relative position in the list of on bits is returned. If bit *i* is off, -1 is returned. Note that legal values of *i* lie on the interval [-2147483648, 2147483647], so the following correction should be made before calling this function:

```
if i > 2147483647: i = i - 4294967296
```

reduceToBitset(*n*)

Returns a ChmBitset of length *n*, which approximates the current ChmSparseBitset. The reduction is done by dividing the coordinate $[0, 2^{32}-1]$ into *n* ranges of equal width, and mapping all sparse bits in a given range to a single bit in the ChmBitset object. For example, if *n* = 1024, the width of each range is $2^{32}/1024 = 4194304$, so sparse bits on the interval $[k \cdot 4194304, (k+1) \cdot 4194304 - 1]$ are mapped to bit *k* in the returned ChmBitset. Thus if any sparse bit in that range is on, bit *k* will be on in the ChmBitset.

simBuser(*bitset*)

Buser similarity.

simCosine(*bitset*)

Cosine similarity.

simDice(*bitset*)

Dice similarity.

simHamann(*bitset*)

Hamann similarity.

simKulczynski(*bitset*)

Kulczynski similarity.

simMatching(*bitset*)

Matching similarity.

`simMcConnaughey(bitset)`
McConnaughey similarity.

`simMinMax(bitset)`
MinMax similarity.

`simModifiedTanimoto(bitset)`
Modified Tanimoto similarity.

`simPearson(bitset)`
Pearson similarity.

`simPetke(bitset)`
Petke similarity.

`simRogersTanimoto(bitset)`
Rogers Tanimoto.

`simSimpson(bitset)`
Simpson similarity.

`simTanimoto(bitset)`
Tanimoto similarity.

`simTversky(bitset, α , β)`
Tversky similarity.

`simYule(bitset)`
Yule similarity.

`size()`
The length of the *bitset*, i.e., 2^{32} .

Fingerprint APIs

The classes and methods in this section support the creation and I/O of binary fingerprints from chemical structures.

Atom Typing Schemes

A Canvas fingerprint is generated by applying a set of rules that decompose a structure into a set of fragments/features, each of which is hashed to a 32-bit integer code that turns on a bit in an underlying `ChmSparseBitset` object. The integer code that's generated depends upon the way in which atoms and bonds are distinguished, and the available schemes are summarized below.

Scheme	Description
1	All atoms and bonds are equivalent.
2	Atoms are distinguished by whether they are hydrogen bond (HB) acceptors or donors; all bonds are equivalent.
3	Atoms are distinguished by hybridization state; all bonds are equivalent.
4	Atoms are distinguished by functional type: {H}, {C}, {F,Cl}, {Br,I}, {N,O}, {S}, {other}. All bonds are equivalent.
5	Sybyl Mol2 atom types; all bonds are equivalent.
6	Atoms are distinguished by whether they are terminal, halogen, HB acceptor/donor; all bonds are equivalent.
7	Atomic number and bond order.
8	Atoms are distinguished by ring size, aromaticity, HB acceptor/donor, whether terminal, whether halogen; bonds are distinguished by bond order.
9	Carhart atom types (atom-pairs approach); all bonds are equivalent.
10	Daylight invariant atom types; bonds are distinguished by bond order.
11	Same as 7, but distinguishing aromatic from non-aromatic.
12	Same as 10, but distinguishing cyclic aliphatic from acyclic aliphatic.

ChmDendriticOut32. Creates and outputs a dendritic fingerprint of theoretical length 2^{32} . Dendritic fingerprints encode linear and branched paths, where a branched path consists of two or more intersecting linear paths.

Constructors

ChmDendriticOut32([*atomTypingScheme*])

The integer *atomTypingScheme* is 10 by default, and it must lie in the range 1-12.

Methods

close()

Closes the binary file on which the object is opened.

close(*minOnFreq*, *maxOffFreq*, *maxBits*[, *reduceBits*])

Filters the bits according to their statistics across all fingerprints in the file, then closes the file. A bit will be eliminated unless the fraction of the time it is on is at least *minOnFreq* and the fraction of the time it is off is no more than *maxOffFreq*. These two parameters work as described only when *maxBits* is 0 and *reduceBits* is 0 or omitted. To keep only the *maxBits* most statistically significant bits, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to the desired value, and set *reduceBits* to 0, or omit it. To reduce the length of the fingerprint

from 2^{32} to $2^{32-\text{reduceBits}}$, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to 0, and set *reduceBits* to the desired value.

`generate(mol)`

Generates a fingerprint from a ChmMol and returns it as a ChmSparseBitset.

`generateMap(mol)`

Generates a fingerprint from a ChmMol, maps the bits back to the atoms and bonds in the structure, and returns the mappings as a ChmFPMMap32.

`getCurrentRowCount()`

The number of fingerprints written to the binary file on which the object is opened.

`getMaxPath()`

The maximum number of bonds in linear paths

`getMinPath()`

The minimum number of bonds in linear paths.

`getTypeInfo()`

A textual summary of the fingerprint settings.

`open(fileName)`

Opens a binary file to which fingerprints may be written.

`setMaxPath(max)`

Sets the maximum number of bonds in linear paths. The default is 5.

`setMinPath(min)`

Sets the minimum number of bonds in linear paths. The default is 0.

`write(bitset, id)`

Writes the fingerprint and a string identifier to the binary file on which the object is opened. Must call the generate method to get *bitset*.

ChmLinearOut32. Creates and outputs a linear fingerprint of theoretical length 2^{32} . Linear fingerprints encode linear paths and ring closure.

Constructors

`ChmLinearOut32([atomTypingScheme])`

The integer *atomTypingScheme* is 10 by default, and it must lie in the range 1-12.

Methods

`close()`

Closes the binary file on which the object is opened.

`close(minOnFreq, maxOffFreq, maxBits[, reduceBits])`

Filters the bits according to their statistics across all fingerprints in the file, then closes the file. A bit will be eliminated unless the fraction of the time it is on is at least *minOnFreq* and the fraction of the time it is off is no more than *maxOffFreq*. These two parameters work as described only when *maxBits* is 0 and *reduceBits* is 0 or omitted. To keep only the *maxBits* most statistically significant bits, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to the desired value, and set *reduceBits* to 0, or omit it. To reduce the length of the fingerprint from 2^{32} to $2^{32-\text{reduceBits}}$, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to 0, and set *reduceBits* to the desired value.

`generate(mol)`

Generates a fingerprint from a ChmMol and returns it as a ChmSparseBitset.

`generateMap(mol)`

Generates a fingerprint from a ChmMol, maps the bits back to the atoms and bonds in the structure, and returns the mappings as a ChmFPMap32.

`getCurrentRowCount()`

The number of fingerprints written to the binary file on which the object is opened.

`getHalfStep()`

True if fragments can end with a bond.

`getMaxPath()`

The maximum number of bonds in linear paths.

`getMaxRingPath()`

The maximum number of bonds that will be traversed in order to close a ring.

`getMinPath()`

The minimum number of bonds in linear paths.

`getTypeInfo()`

A textual summary of the fingerprint settings.

`open(fileName)`

Opens a binary file to which fingerprints may be written.

`setHalfStep(doHalf)`

Controls whether linear fragments can start/end at a bond. If False, all linear fragments will start and end at an atom, which is the default treatment.

`setMaxPath(max)`

Sets the maximum number of bonds in linear paths. The default is 7.

`setMaxRingPath(max)`

Sets the maximum number of bonds that will be traversed in order to close a ring. The default is 14. If 0, all paths encoded in the fingerprint will be self-avoiding.

`setMinPath(min)`

Sets the minimum number of bonds in linear paths. The default is 0.

`write(bitset, id)`

Writes the fingerprint and a string identifier to the binary file on which the object is opened. Must call the generate method to get *bitset*.

ChmMolprint2D32. Creates and outputs MOLPRINT 2D fingerprints of theoretical length 2^{32} . MOLPRINT 2D fingerprints encode atom environments using lists of atom types located at topological distances of 0 to 2 bonds from each heavy atom.

Constructors

`ChmMolprint2D32([atomTypingScheme])`

The integer *atomTypingScheme* is 5 by default, and it must lie in the range 1-12.

Methods

`close()`

Closes the binary file on which the object is opened.

`close(minOnFreq, maxOffFreq, maxBits[], reduceBits)`

Filters the bits according to their statistics across all fingerprints in the file, then closes the file. A bit will be eliminated unless the fraction of the time it is on is at least *minOnFreq* and the fraction of the time it is off is no more than *maxOffFreq*. These two parameters work as described only when *maxBits* is 0 and *reduceBits* is 0 or omitted. To keep only the *maxBits* most statistically significant bits, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to the desired value, and set *reduceBits* to 0, or omit it. To reduce the length of the fingerprint from 2^{32} to $2^{32-\text{reduceBits}}$, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to 0, and set *reduceBits* to the desired value.

`generate(mol)`

Generates a fingerprint from a ChmMol and returns it as a ChmSparseBitset.

`generateMap(mol)`

Generates a fingerprint from a ChmMol, maps the bits back to the atoms and bonds in the structure, and returns the mappings as a ChmFPMMap32.

`getCurrentRowCount()`

The number of fingerprints written to the binary file on which the object is opened.

getTypeInfo()

A textual summary of the fingerprint settings.

open(fileName)

Opens a binary file to which fingerprints may be written.

write(bitset, id)

Writes the fingerprint and a string identifier to the binary file on which the object is opened. Must call the generate method to get *bitset*.

ChmPairwiseOut32. Creates and outputs a pairwise fingerprint of theoretical length 2^{32} . Pairwise fingerprints encode pairs of atoms, differentiated by type, and the distance separating them: $Type_i - Type_j - d_{ij}$.

Constructors

ChmPairwiseOut32([atomTypingScheme])

The integer *atomTypingScheme* is 9 by default, and it must lie in the range 1-12.

Methods

close()

Closes the binary file on which the object is opened.

close(minOnFreq, maxOffFreq, maxBits[, reduceBits])

Filters the bits according to their statistics across all fingerprints in the file, then closes the file. A bit will be eliminated unless the fraction of the time it is on is at least *minOnFreq* and the fraction of the time it is off is no more than *maxOffFreq*. These two parameters work as described only when *maxBits* is 0 and *reduceBits* is 0 or omitted. To keep only the *maxBits* most statistically significant bits, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to the desired value, and set *reduceBits* to 0, or omit it. To reduce the length of the fingerprint from 2^{32} to $2^{32-reduceBits}$, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to 0, and set *reduceBits* to the desired value.

generate(mol)

Generates a fingerprint from a ChmMol and returns it as a ChmSparseBitset.

generateMap(mol)

Generates a fingerprint from a ChmMol, maps the bits back to the atoms and bonds in the structure, and returns the mappings as a ChmFPMMap32.

getCurrentRowCount()

The number of fingerprints written to the binary file on which the object is opened.

`getTypeInfo()`

A textual summary of the fingerprint settings.

`open(fileName)`

Opens a binary file to which fingerprints may be written.

`set3D(use3D)`

Controls whether 3D distances will be used. If False, all distances will be 2D, which is the default treatment. If True, `setBinWidth` must also be called.

`setBinWidth(dbin)`

Sets the bin width for 2D/3D distances. The default for 2D distances is 1. There is no default for 3D distances, so this function must be called if using 3D distances.

`setFuzzyDistance(dfuzz)`

If a 2D or 3D distance is within *dfuzz* of a boundary, turn on the bit corresponding to the neighboring bin.

`setMaxDistance(dmax)`

Sets the maximum 2D/3D distance to consider. The default is no limit.

`setMinDistance(dmin)`

Sets the minimum 2D/3D distance to consider. The default is 0.

`write(bitset, id)`

Writes the fingerprint and a string identifier to the binary file on which the object is opened. Must call the `generate` method to get *bitset*.

ChmRadialOut32. Creates and outputs a radial fingerprint of theoretical length 2^{32} . Radial fingerprints (also known as extended connectivity fingerprints) encode fragments that grow radially from each heavy atom.

Constructors

`ChmRadialOut32([atomTypingScheme])`

The integer *atomTypingScheme* is 4 by default, and it must lie in the range 1-12.

Methods

`close()`

Closes the binary file on which the object is opened.

`close(minOnFreq, maxOffFreq, maxBits[, reduceBits])`

Filters the bits according to their statistics across all fingerprints in the file, then closes the file. A bit will be eliminated unless the fraction of the time it is on is at least *minOnFreq* and the fraction of the time it is off is no more than *maxOffFreq*. These two parameters work as described only when *maxBits* is 0 and *reduceBits* is 0 or omitted. To keep only the *maxBits* most statistically significant bits, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to the desired value, and set *reduceBits* to 0, or omit it. To reduce the length of the fingerprint from 2^{32} to $2^{32-\text{reduceBits}}$, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to 0, and set *reduceBits* to the desired value.

`generate(mol)`

Generates a fingerprint from a ChmMol and returns it as a ChmSparseBitset.

`generateMap(mol)`

Generates a fingerprint from a ChmMol, maps the bits back to the atoms and bonds in the structure, and returns the mappings as a ChmFPMMap32.

`getCurrentRowCount()`

The number of fingerprints written to the binary file on which the object is opened.

`getIterations()`

The number of iterations of radial growth, which is also the radial size of the largest fragments.

`getMinSize()`

The minimum radial size of fragments that will be retained.

`getTypeInfo()`

A textual summary of the fingerprint settings.

`open(fileName)`

Opens a binary file to which fingerprints may be written.

`setIterations(iter)`

Sets the number of iterations of radial growth, which is also the radial size of the largest fragments in the fingerprint. The default is 4.

`setMinSize(min)`

Sets the minimum radial size of fragments to retain. All fragments generated prior to iteration *min* will be discarded. The default is 0.

`write(bitset, id)`

Writes the fingerprint and a string identifier to the binary file on which the object is opened. Must call the generate method to get *bitset*.

ChmSMARTSOut32. Creates and outputs a custom fingerprint whose bits are defined by a series of SMARTS patterns in a text file:

SMARTS1 bitName1
SMARTS2 bitName2
etc.

Each SMARTS pattern and bit name must be separated by one or more spaces.

Constructors

`ChmSMARTSOut32(smartsFile[, wantCSV])`

The Boolean *wantCSV* is False by default and determines whether the output fingerprint file should be a CSV file.

Methods

`close()`

Closes the binary/CSV file on which the object is opened.

`close(minOnFreq, maxOffFreq, maxBits[, reduceBits])`

Filters the bits according to their statistics across all fingerprints in the file, then closes the file. A bit will be eliminated unless the fraction of the time it is on is at least *minOnFreq* and the fraction of the time it is off is no more than *maxOffFreq*. These two parameters work as described only when *maxBits* is 0 and *reduceBits* is 0 or omitted. To keep only the *maxBits* most statistically significant bits, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to the desired value, and set *reduceBits* to 0, or omit it. To reduce the length of the fingerprint from 2^{32} to $2^{32-\text{reduceBits}}$, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to 0, and set *reduceBits* to the desired value.

`generate(mol)`

Generates a fingerprint from a ChmMol and returns it as a ChmSparseBitset.

`generateMap(mol)`

Generates a fingerprint from a ChmMol, maps the bits back to the atoms and bonds in the structure, and returns the mappings as a ChmFPMMap32.

`getCurrentRowCount()`

The number of fingerprints written to the binary/CSV file on which the object is opened.

`getTypeInfo()`

A textual summary of the fingerprint settings.

`open(fileName)`

Opens a binary/CSV file to which fingerprints may be written.

`write(bitset, id)`

Writes the fingerprint and a string identifier to the binary/CSV file on which the object is opened. Must call the generate method to get *bitset*.

ChmTopologicalTorsionOut32. Creates and outputs a topological torsion fingerprint of theoretical length 2^{32} . Topological torsion fingerprints encode linear paths of 4 atoms, differentiated by type: $Type_i-Type_j-Type_k-Type_l$.

Constructors

`ChmTopologicalTorsionOut32([atomTypingScheme[, span]])`

The integer *atomTypingScheme* is 10 by default, and it must lie in the range 1-12.

Methods

`close()`

Closes the binary file on which the object is opened.

`close(minOnFreq, maxOffFreq, maxBits[, reduceBits])`

Filters the bits according to their statistics across all fingerprints in the file, then closes the file. A bit will be eliminated unless the fraction of the time it is on is at least *minOnFreq* and the fraction of the time it is off is no more than *maxOffFreq*. These two parameters work as described only when *maxBits* is 0 and *reduceBits* is 0 or omitted. To keep only the *maxBits* most statistically significant bits, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to the desired value, and set *reduceBits* to 0, or omit it. To reduce the length of the fingerprint from 2^{32} to $2^{32-reduceBits}$, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to 0, and set *reduceBits* to the desired value.

`generate(mol)`

Generates a fingerprint from a ChmMol and returns it as a ChmSparseBitset.

`generateMap(mol)`

Generates a fingerprint from a ChmMol, maps the bits back to the atoms and bonds in the structure, and returns the mappings as a ChmFPMap32.

`getCurrentRowCount()`

The number of fingerprints written to the binary file on which the object is opened.

`getTypeInfo()`

A textual summary of the fingerprint settings.

`open(fileName)`

Opens a binary file to which fingerprints may be written.

`write(bitset, id)`

Writes the fingerprint and a string identifier to the binary file on which the object is opened. Must call the generate method to get *bitset*.

ChmTripletOut32. Creates and outputs a triplet fingerprint of theoretical length 2^{32} . Triplet fingerprints encode triplets of atoms, differentiated by type, and the distances separating them: $Type_i-d_{ij}-Type_j-d_{jk}-Type_k-d_{ki}$.

Constructors

ChmTripletOut32([*atomTypingScheme*])

The integer *atomTypingScheme* is 10 by default, and it must lie in the range 1-12.

Methods

close()

Closes the binary file on which the object is opened.

close(*minOnFreq*, *maxOffFreq*, *maxBits*[, *reduceBits*])

Filters the bits according to their statistics across all fingerprints in the file, then closes the file. A bit will be eliminated unless the fraction of the time it is on is at least *minOnFreq* and the fraction of the time it is off is no more than *maxOffFreq*. These two parameters work as described only when *maxBits* is 0 and *reduceBits* is 0 or omitted. To keep only the *maxBits* most statistically significant bits, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to the desired value, and set *reduceBits* to 0, or omit it. To reduce the length of the fingerprint from 2^{32} to $2^{32-\text{reduceBits}}$, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to 0, and set *reduceBits* to the desired value.

generate(*mol*)

Generates a fingerprint from a ChmMol and returns it as a ChmSparseBitset.

generateMap(*mol*)

Generates a fingerprint from a ChmMol, maps the bits back to the atoms and bonds in the structure, and returns the mappings as a ChmFPMap32.

getCurrentRowCount()

The number of fingerprints written to the binary file on which the object is opened.

getTypeInfo()

A textual summary of the fingerprint settings.

open(*fileName*)

Opens a binary file to which fingerprints may be written.

set3D(*use3D*)

Controls whether 3D distances will be used. If False, all distances will be 2D, which is the default treatment. If True, `setBinWidth` must also be called.

`setBinWidth(dbin)`

Sets the bin width for 2D/3D distances. There is no default for 3D distances, so this function must be called if using 3D distances.

`setFuzzyDistance(dfuzz)`

If a 2D or 3D distance is within *dfuzz* of a boundary, turn on the bit corresponding to the neighboring bin.

`setMaxDistance(dmax)`

Sets the maximum 2D/3D distance to consider. The default is no limit.

`setMinDistance(dmin)`

Sets the minimum 2D/3D distance to consider. The default is 0.

`write(bitset, id)`

Writes the fingerprint and a string identifier to the binary file on which the object is opened. Must call the `generate` method to get *bitset*.

ChmFPIn32. Reads a binary fingerprint file created by `ChmDendriticOut32`, `ChmLinearOut32`, etc., and extracts each fingerprint as a `ChmSparseBitset`. A typical workflow to read a file is as follows:

```
fpIn = canvas.ChmFpIn("linear.fp")
while fpIn.hasNext():
    fp = fpIn.next # A ChmSparseBitset
```

Constructors

`ChmFPIn32(fileName)`

Creates the object and opens a binary fingerprint file.

Methods.

`filter(minOnFreq, maxOffFreq, maxBits, reduceBits)`

Filters the bits according to their statistics across all fingerprints in the file, then rewrites the file. A bit will be eliminated unless the fraction of the time it is on is at least *minOnFreq* and the fraction of the time it is off is no more than *maxOffFreq*. These two parameters work as described only when *maxBits* is 0 and *reduceBits* is 0. To keep only the *maxBits* most statistically significant bits, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to the desired value, and set *reduceBits* to 0. To reduce the length of the fingerprint from 2^{32} to $2^{32-\text{reduceBits}}$, set *minOnFreq* to 0.0, set *maxOffFreq* to 1.0, set *maxBits* to 0, and set *reduceBits* to the desired value. If *reduceBits* is applied to the same file repeatedly, the fingerprint will become progressively shorter, until it contains only one bit.

getFileName()

The name of the binary fingerprint file on which the object is opened.

getOnBits()

A ChmSparseBitset that contains the union of on bits in the fingerprint file.

getPos()

The number of fingerprints that have been read from the file on which the object is opened.

getRowCount()

The number of fingerprints contained in the file on which the object is opened.

getTypeInfo()

A textual summary of the fingerprint settings.

getUnionCount()

Equivalent to getOnBits().count().

hasNext ()

True if more fingerprints remain in the file.

next ()

The next fingerprint in the file, as a ChmSparseBitset of length 2^{32} .

nextReduced ()

The next fingerprint in the file, as a ChmBitset of length getUnionCount().

rewind ()

Rewinds the file to the beginning.

ChmFPMMap32. Provides a mapping of each fingerprint bit back to the atoms and bonds that are responsible for setting that bit. A ChmFPMMap32 object is obtained by calling the generateMap method of ChmDendriticOut32, ChmLinearOut32, etc.

Methods

getCode(i)

The 32-bit integer code for the i^{th} on bit, where i runs from 0 to getCodeCount()-1. This is the position of the on bit in an underlying ChmSparseBitset object. Note that if the underlying unsigned value is larger than 2147483647, a negative value will be returned. In other words, if $\text{code} > 2147483647$, $\text{code} \rightarrow \text{code} - 4294967296$. The returned value can be safely used in all subsequent APIs that accept a code.

getCodeCount()

The number of unique codes in the fingerprint, i.e., the number of on bits.

`getAtomsMaskForCode(code)`

Iterable list of `ChmBitset` objects, each of which encodes a set of atoms that are associated with the bit. If the length of the list is greater than one, the bit maps to more than one set of atoms.

`getBondsMaskForCode(code)`

Iterable list of `ChmBitset` objects, each of which encodes a set of bonds that are associated with the bit. If the length of the list is greater than one, the bit maps to more than one set of bonds.

`getOnBits()`

Returns the underlying `ChmSparseBitset` object that holds the fingerprint.

`hasCode(code)`

True if *code* is the position of an on bit.

ChmPairwiseFPIn. Reads and stores a matrix of similarities supplied in binary or CSV format.

Matrix File Format Constants

These constants are scoped at the canvas module level, so the definitions below assume that the module has been imported as follows:

```
import schrodinger.application.canvas.base as canvas
```

`canvas.Binary` – Binary format (0).

`canvas.CSV` – Comma-separated values, with column and row headers (1).

`canvas.CSV_NoColHeader` – CSV file does not contain column names on first line (2).

`canvas.CSV_NoRowHeader` – CSV file does not contain row names in the first column (3).

`canvas.CSV_NoRowOrColHeader` – CSV file does not contain row names or column names (4).

Constructors

`ChmPairwiseFPIn(inFile, matrixFormat[, forceSquare])`

Reads a similarity/distance matrix from the file *inFile*. The integer *matrixFormat* must be one of the values in “Matrix File Format Constants.” The Boolean *forceSquare* is False by default and controls whether the matrix must be square.

`ChmPairwiseFPIn(inStream, matrixFormat[, forceSquare])`

Reads a similarity/distance matrix from an input stream. Use `canvas ifstream(inFile)` to create an input stream on a file named *inFile*.

Methods

`cols()`

The number of columns in the matrix.

`getColumnName(i)`

The i^{th} column name. *i* runs from 0 to `cols()-1`.

`getRowName(i)`

The i^{th} row name. *i* runs from 0 to `rows()-1`.

`isDistanceMatrix()`

True if the matrix is square and the diagonals are 0.

`isSimilarityMatrix()`

True if the matrix is square and the diagonals are 1.

`isSquare()`

True if the matrix is square.

`isSymmetric()`

True if the matrix is symmetric.

`isValid()`

True if the matrix is square and the diagonals are 1 or 0.

`rows()`

The number of rows in the matrix.

ChmPairwiseFPOut32. Given one or two fingerprint files created by `ChmDendriticOut32`, `ChmLinearOut32`, etc. this class computes and outputs a similarity or distance matrix in binary or CSV format. Not to be confused with `ChmPairwiseOut32`, which computes and outputs a pairwise fingerprint.

Similarity/Distance Method Constants

One of these integer codes must be supplied when calling the `setMetric` method.

- 1 – Buser similarity
- 2 – Cosine similarity
- 3 – Dice similarity
- 4 – Dixon distance

- 5 – Euclidean distance
- 6 – Hamann similarity
- 7 – Hamming distance
- 8 – Kulczynski similarity
- 9 – Matching similarity
- 10 – McConnaughey similarity
- 11 – MinMax
- 12 – Modified Tanimoto
- 13 – Pattern difference distance
- 14 – Pearson similarity
- 15 – Petke similarity
- 16 – Rogers Tanimoto similarity
- 17 – Shape distance
- 18 – Simpson similarity
- 19 – Size distance
- 20 – Soergel distance
- 21 – Tanimoto similarity (default)
- 22 – Tversky similarity
- 23 – Variance similarity
- 24 – Yule similarity

Constructors

`ChmPairwiseFPOut32(outFile, forceBinary, α , β)`

The matrix file *outFile* is binary by default, but it may be changed to CSV by calling `setHumanReadable(True)`. The Boolean *forceBinary* controls whether fingerprints with scaled bit values are treated as ordinary 0/1 values. This parameter is relevant only when processing fingerprints created `canvasFPGen` with the `–scaling` option, and it should normally be set to `True`. The floating point parameters α and β are applied only when Tversky similarities are computed, and they may be set to 0 otherwise.

Methods

`getFlattening()`

True if similarities will be flattened through an exponential transformation.

`getHumanReadable()`

True if the matrix file will be written in CSV format.

`getMetric()`

The integer-valued similarity/distance method.

`process(fpFile[, blockSize])`

Computes and outputs a symmetric matrix of self-similarities/distances for the fingerprints in *fpFile*, which must be a binary file. The integer *blockSize* is 100 by default, and it determines how many fingerprints will be held in memory at one time.

`process(fpFile, start, stop[, blockSize])`

Computes and outputs a symmetric matrix of self-similarities/distances for a subset of the fingerprints in *fpFile*. The integers *start* and *stop* are zero-based indices. If *stop* is greater than or equal to the number of fingerprints in the file, the matrix will be padded with zeros.

`process(fpFile1, fpFile2[, blockSize])`

Computes and outputs a rectangular matrix of similarities/distances for two fingerprint files. Rows are spanned by *fpFile1* and columns are spanned by *fpFile2*. The two files must contain the same type of fingerprint.

`process(fpFile1, start1, stop1, fpFile2, start2, stop2[, blockSize])`

Computes and outputs a rectangular matrix of similarities/distances for subsets of fingerprints in two files.

`setFlattening(coeff)`

Subject similarities to the following exponential transformation that drives smaller values toward zero: $sim \rightarrow \exp[-coeff \cdot (1-sim)^2]$. A reasonable value of *coeff* is 25. Not applied when computing distances.

`setHumanReadable(writeCSV)`

The Boolean *writeCSV* controls whether similarities/distances will be written in CSV format. The default format is binary.

`setMetric(code)`

The integer *code* controls how similarities/distances will be computed and it must lie in the range 1-24. The default is 21 (Tanimoto similarity).

Hierarchical Clustering APIs

The classes and methods in this section support hierarchical agglomerative clustering on a matrix of similarities or distances.

ChmHierarchicalClustering. Performs a full hierarchical, agglomerative clustering procedure, merging *n* objects into *n*−1 clusters, *n*−2 clusters,..., 1 cluster.

Linkage Method Constants

Determines how the distance between two clusters is computed.

ChmHierarchicalClustering.SingleLinkage – Closest inter-cluster pair (1).

ChmHierarchicalClustering.CompleteLinkage – Farthest inter-cluster pair (2).

ChmHierarchicalClustering.GroupAverageLinkage – Average distance between all inter-cluster pairs (3).

ChmHierarchicalClustering.UnweightedCentroidLinkage – Euclidean distance between cluster centroids (4).

ChmHierarchicalClustering.WeightedAverageLinkage – Average distance to the two clusters merged in forming a given cluster (5).

ChmHierarchicalClustering.WardsMinimumVarianceLinkage – Sum of squared distances to merged cluster centroid (6).

ChmHierarchicalClustering.WeightedCentroidLinkage – Weighted center of mass distance, aka median (7).

ChmHierarchicalClustering.FlexibleBetaLinkage – Weighted average intra-cluster and inter-cluster distances, aka Lance-Williams (8).

ChmHierarchicalClustering.SchrodingerLinkage – Closest distance between terminal ($R \leftrightarrow L$) points in 1D cluster orderings (9).

Pairwise Selection Constants

Indicates the preference for using the upper or lower triangle of the similarity/distance matrix M in case it's not symmetric.

ChmHierarchicalClustering.UseEither – No preference; matrix is symmetric (0).

ChmHierarchicalClustering.UseMinAB – Use $\min\{M(i, j), M(j, i)\}$ (1).

ChmHierarchicalClustering.UseMaxAB – Use $\max\{M(i, j), M(j, i)\}$ (2).

ChmHierarchicalClustering.UseUpper – Use $M(i, j)$, where $i < j$ (3).

ChmHierarchicalClustering.UseLower – Use $M(i, j)$, where $i > j$ (4).

Constructors

ChmHierarchicalClustering([linkageMethod][, autoScale])

The integer *linkageMethod* is ChmHierarchicalClustering.GroupAverageLinkage by default, and the Boolean *autoScale* is True by default. *autoScale* affects only the stress value reported by the cluster methods, not how the actual clustering is done.

Methods

cluster(pairwiseIn, treeStreamOut, selectionType)

Performs clustering on a *ChmPairwiseIn* object and returns a value that measures the level of stress in the 1D ordering. Information required to create a dendrogram is written to the output stream *treeStreamOut*, which can be created using `canvas.ofstream(treeFile)`, where *treeFile* is name of the destination file. *treeFile* can be supplied to the canvas utility `canvas.TreeDraw` to create a PostScript dendrogram. *selectionType* must be one of the values in “Pairwise Selection Constants.”

`cluster(pairwiseIn, treeStreamOut, baseStreamOut, selectionType)`

Writes a summary of the clustering process to the output stream *baseStreamOut*. See *ChmClusterGrouper* for subsequent usage of the *baseFile* associated with *baseStreamOut*. The return value and other parameters are as described previously.

`cluster(pairwiseIn, treeStreamOut, baseStreamOut, statsStreamOut, selectionType, allowTies)`

Writes statistics at each clustering level to the CSV output stream *statsStreamOut*. The Boolean *allowTies* controls whether more than one cluster member can be flagged as nearest to centroid or farthest from centroid if those members are equidistant to the centroid. This parameter has no effect on clustering and is actually relevant only when groupings are output (see next method). The return value and other parameters are as described previously.

`cluster(pairwiseIn, treeStreamOut, groupStreamOut, n, selectionType, noSingletons, allowTies)`

Writes groupings for the formation of *n* clusters to the CSV output stream *groupStreamOut*. If the Boolean *noSingletons* is True, groupings will be reported only if there are at least *n* non-singleton clusters. Otherwise, *groupStreamOut* will contain only singletons (i.e., no groupings). The return value and other parameters are as described previously.

`cluster(pairwiseIn, treeStreamOut, groupStreamOut, d, selectionType, noSingletons, allowTies)`

Directly analogous to the previous method except that *d* is a threshold on the merging distance. All clusters formed at or below this distance will be written to *groupStreamOut*.

ChmClusterGrouper. Generates cluster groupings from the *baseFile* created by a *ChmHierarchicalClustering* object, or by the utilities `canvasHC` and `canvasHCBuild`.

Constructors

`ChmClusterGrouper()`

Instantiates a *ChmClusterGrouper* object.

Methods

`group(pairwiseIn, baseStreamIn, groupStreamOut, n, noSingletons, allowTies)`

Given a *ChmPairwiseFIn* object *pairwiseIn* and an associated input stream *baseStreamIn* that's opened on a *baseFile*, this function writes groupings for the formation of *n* clusters to the CSV output stream *groupStreamOut*. If the Boolean *noSingletons* is True, groupings will be reported only if there are at least *n* non-singleton clusters. Otherwise, *groupStreamOut*

will contain only singletons (i.e., no groupings). The Boolean *allowTies* controls whether more than one cluster member can be flagged as nearest to centroid or farthest from centroid if those members are equidistant to the centroid.

`groupByThreshold(pairwiseIn, baseStreamIn, groupStreamOut, d, allowTies)`

Analogous to the previous method except that d is a threshold on the merging distance. All clusters formed at or below this distance will be written to *groupStreamOut*.

`getDforN(pairwiseIn, baseStreamIn, n)`

Returns the merging distance for the formation of n clusters.

`getNforD(pairwiseIn, baseStreamIn, d)`

Returns the number of clusters formed at below the merging distance d .